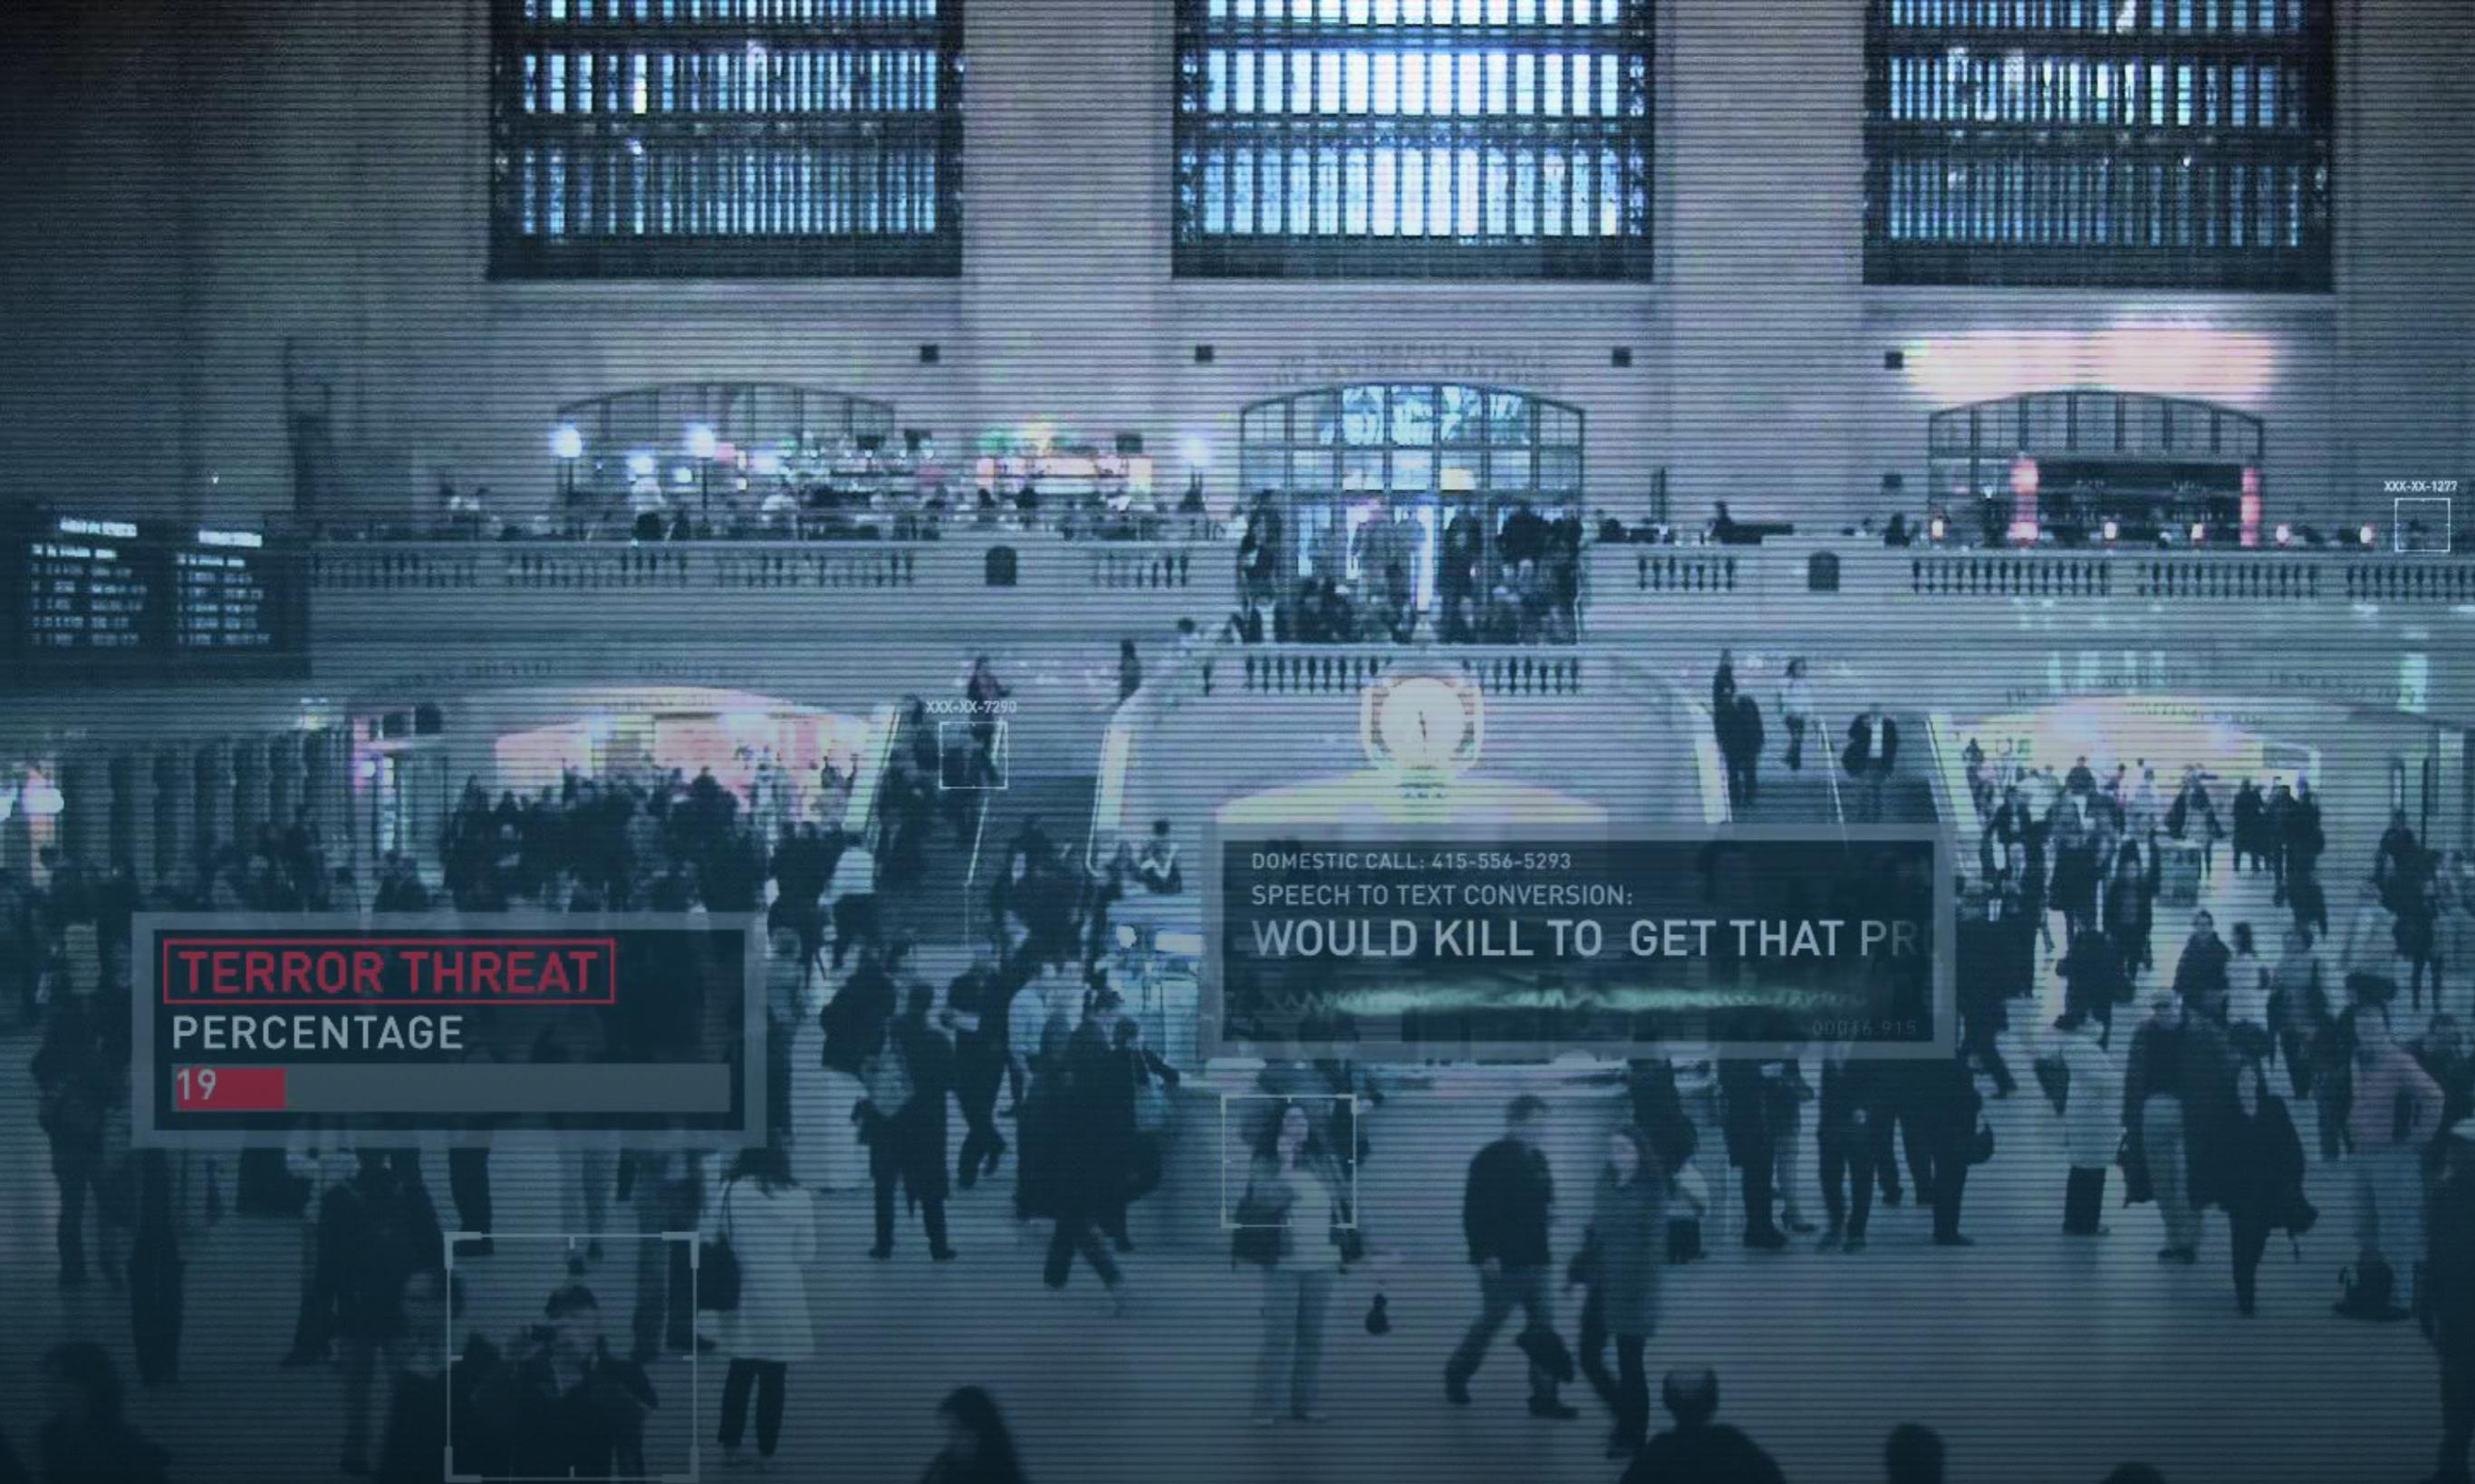




DevFest 2017

SIDDHARTH KOTHIYAL • AVIKANT SAINI

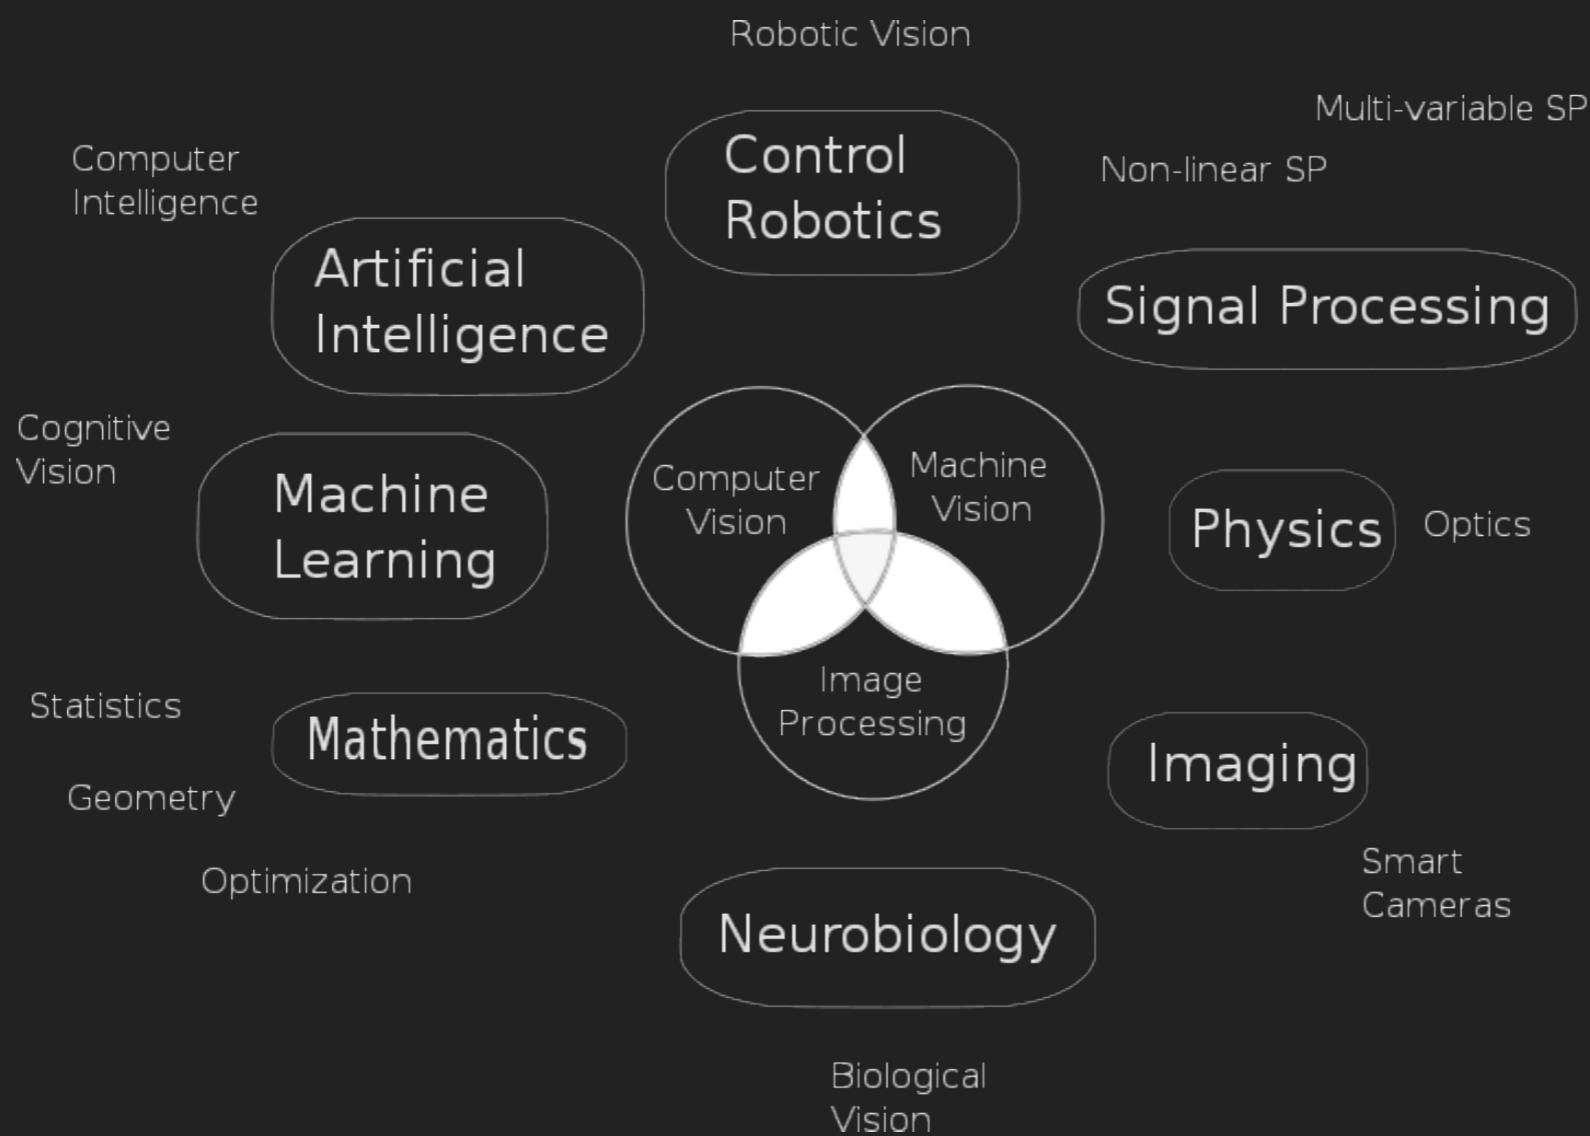
IMPLEMENTATION OF COMPUTER VISION ALGORITHMS



WHAT IS COMPUTER VISION?

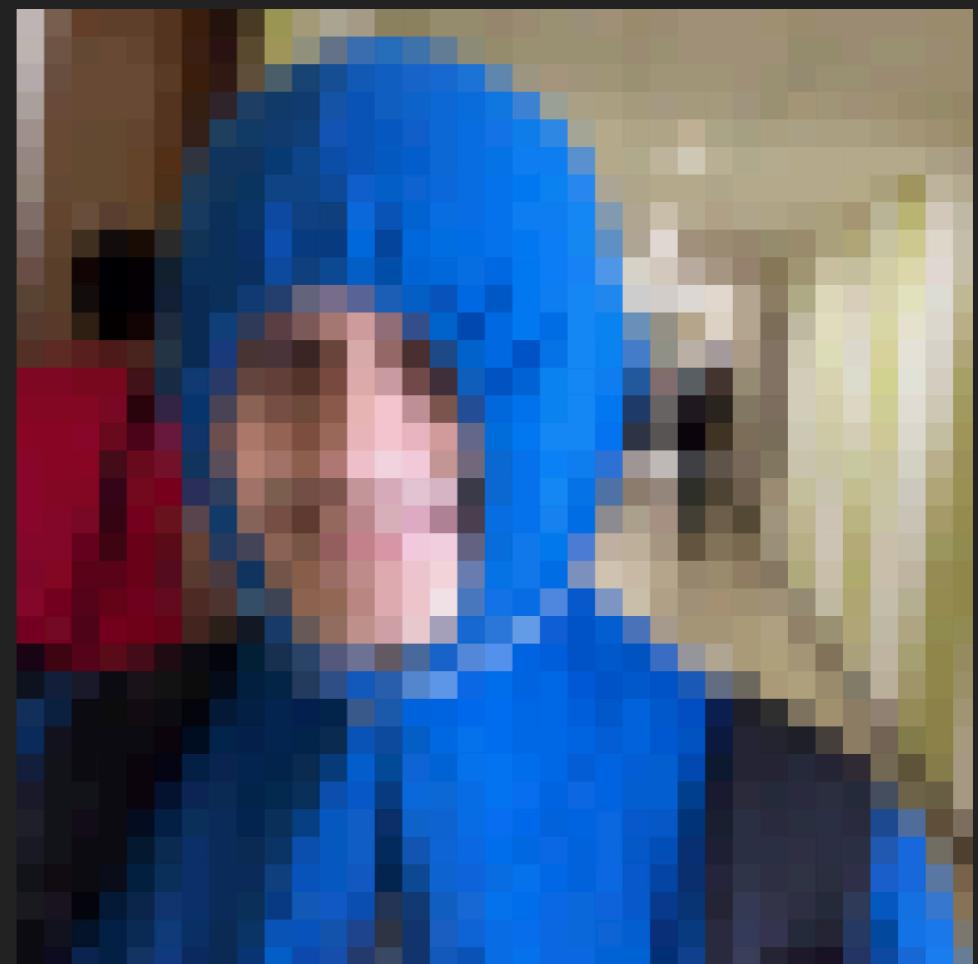
WHAT IS COMPUTER VISION

- Image → Sensor → Interpreter → Information



GOAL OF COMPUTER VISION

- ▶ Extract meaning from pixels
- ▶ Humans be good at this
- ▶ Computers not so much

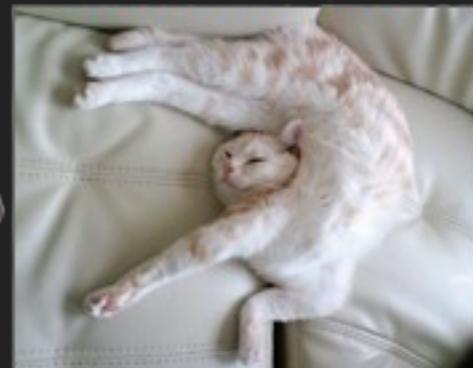


KINDS OF COMPUTER VISION

- ▶ LOW LEVEL
 - ▶ Measurements
 - ▶ Enhancements
 - ▶ Region segmentation
 - ▶ Features
- ▶ MID LEVEL
 - ▶ Reconstruction
 - ▶ Depth
 - ▶ Motion Estimation
- ▶ HIGH LEVEL
 - ▶ Category Detection
 - ▶ Activity Recognition
 - ▶ Deep Understanding

WHY BE COMPUTER VISION DIFFICULT

- ▶ Viewpoint variation
- ▶ Deformation
- ▶ Motion
- ▶ Illumination changes
- ▶ Occlusion
- ▶ Object class variation
- ▶ Scale
- ▶ Background Clutter
- ▶ Local ambiguity



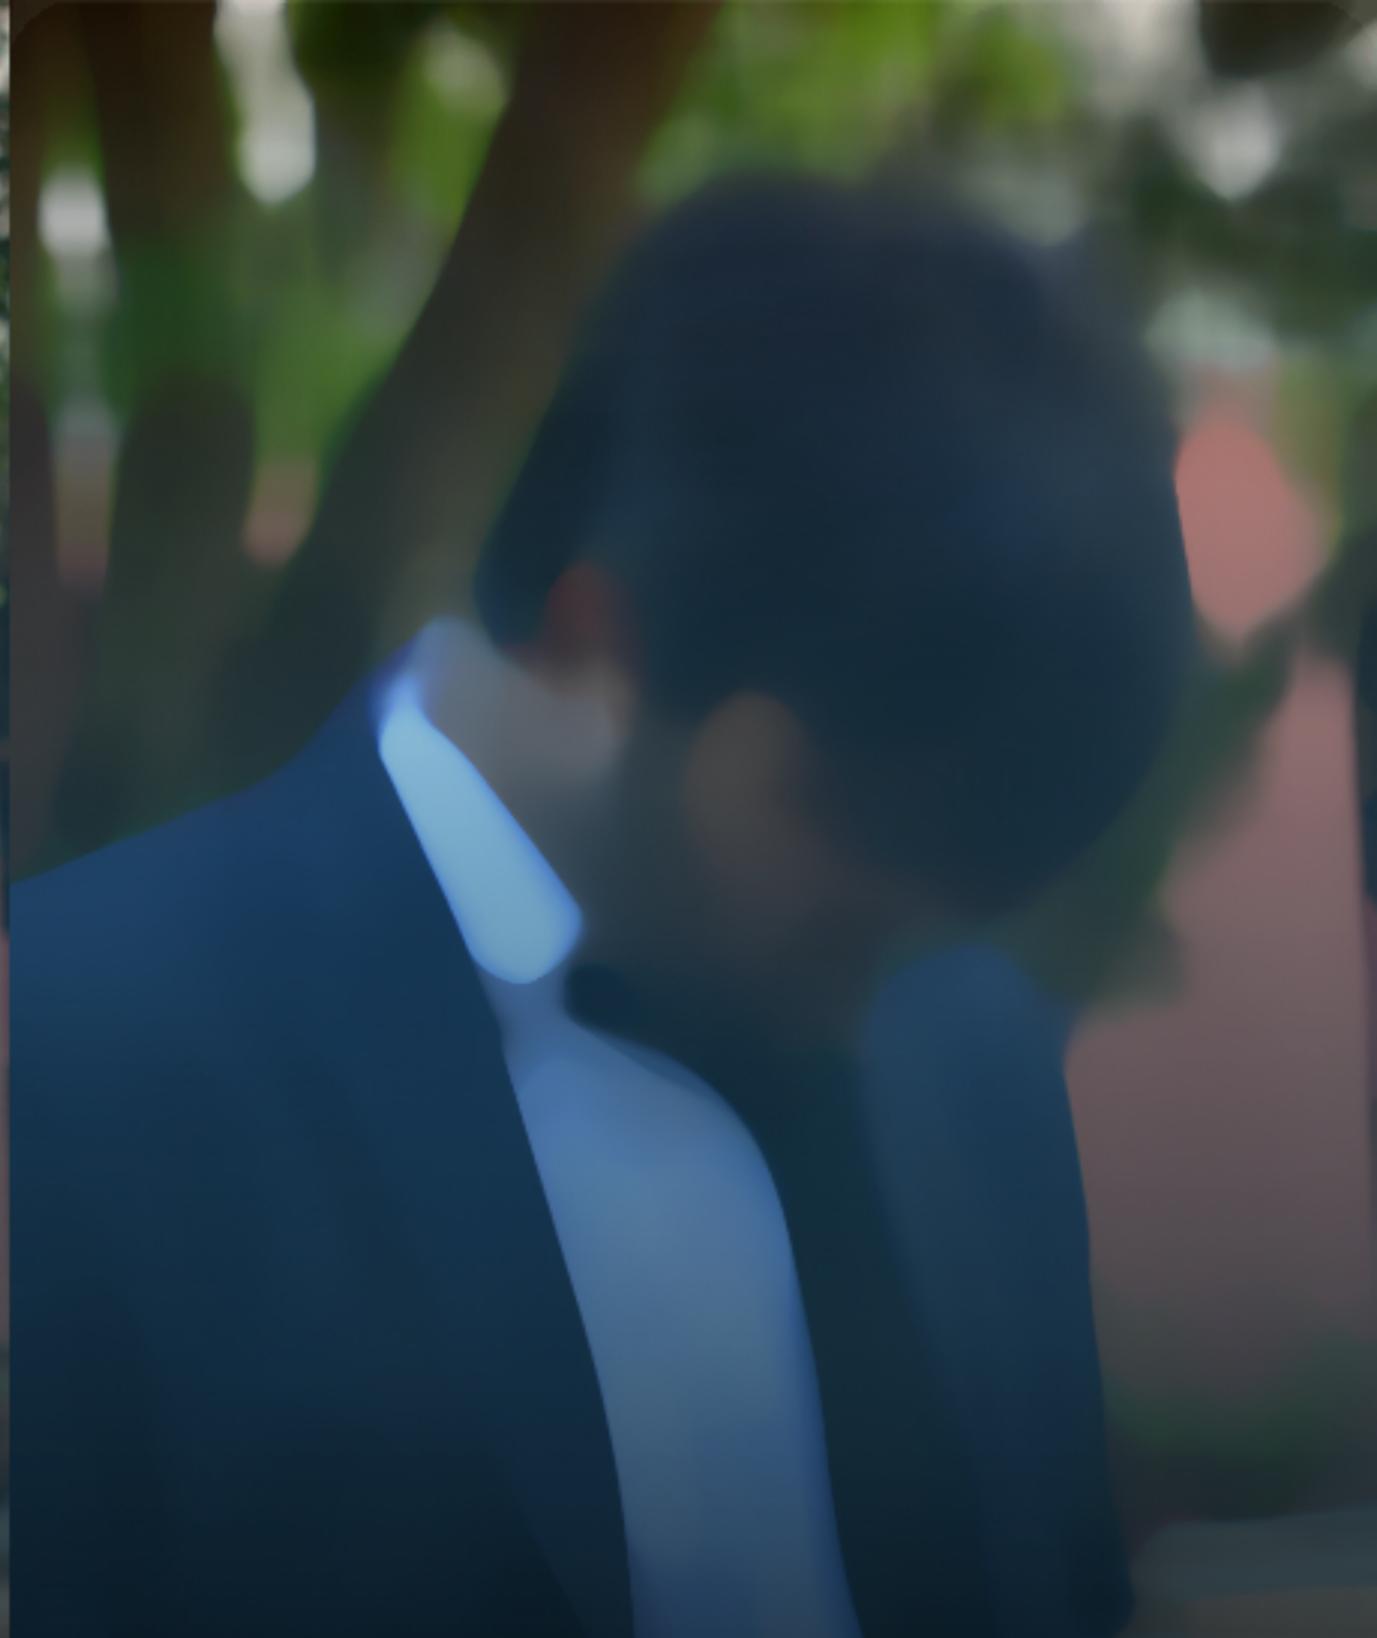
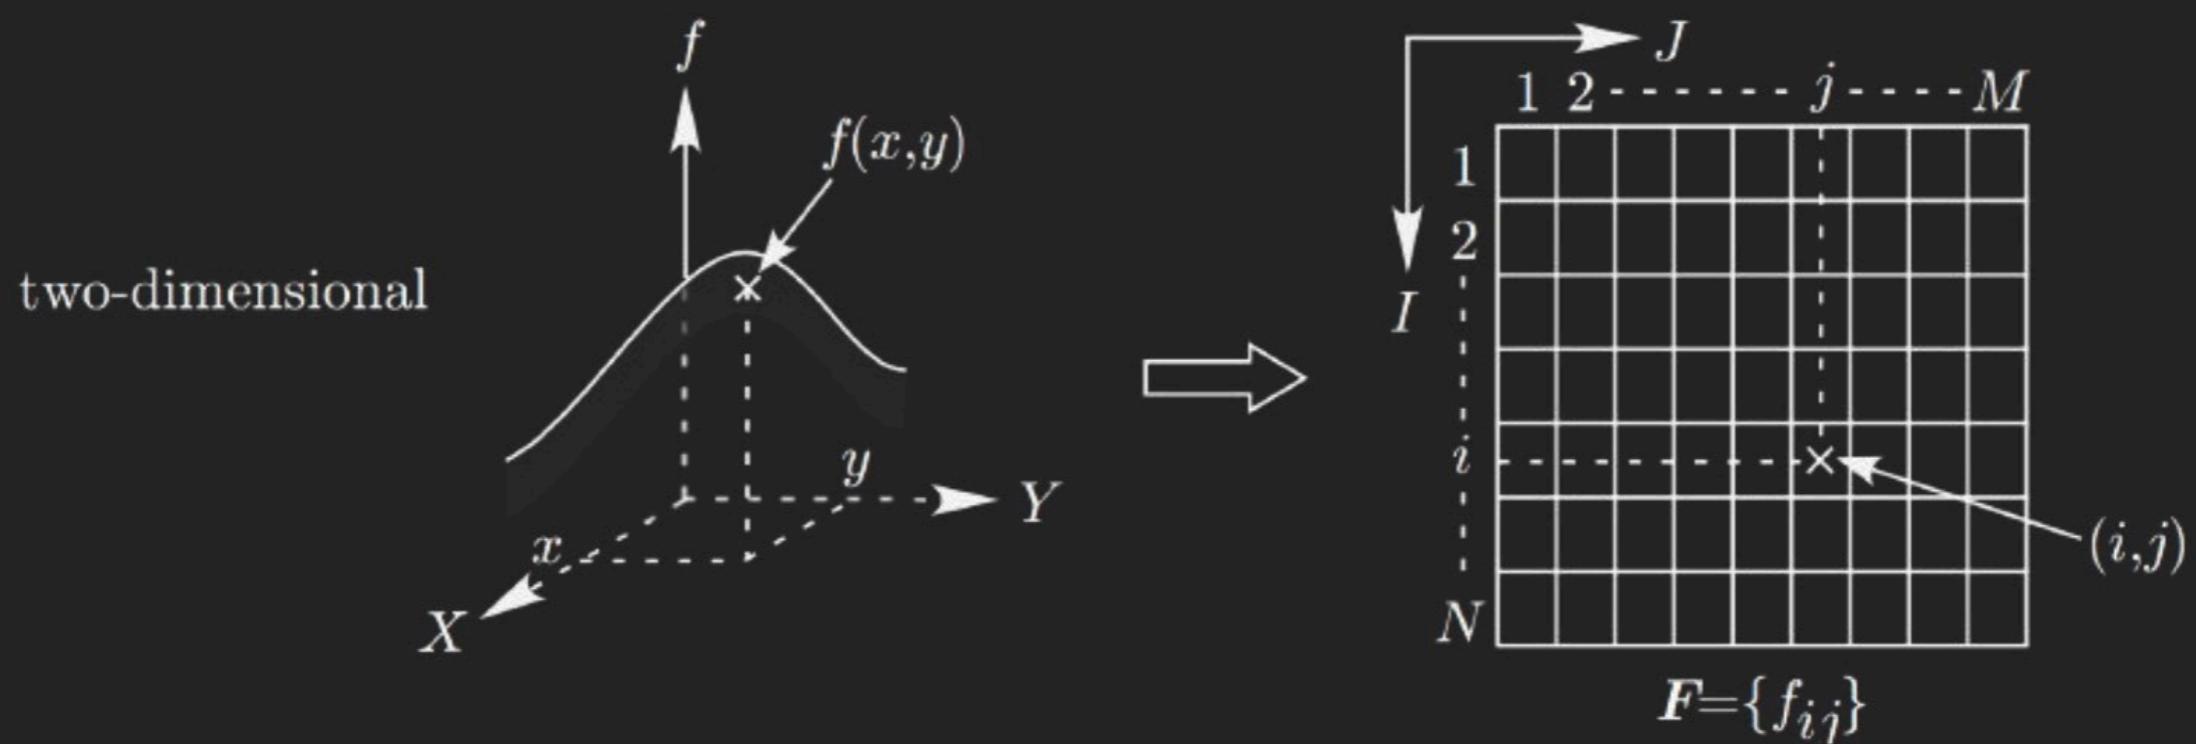
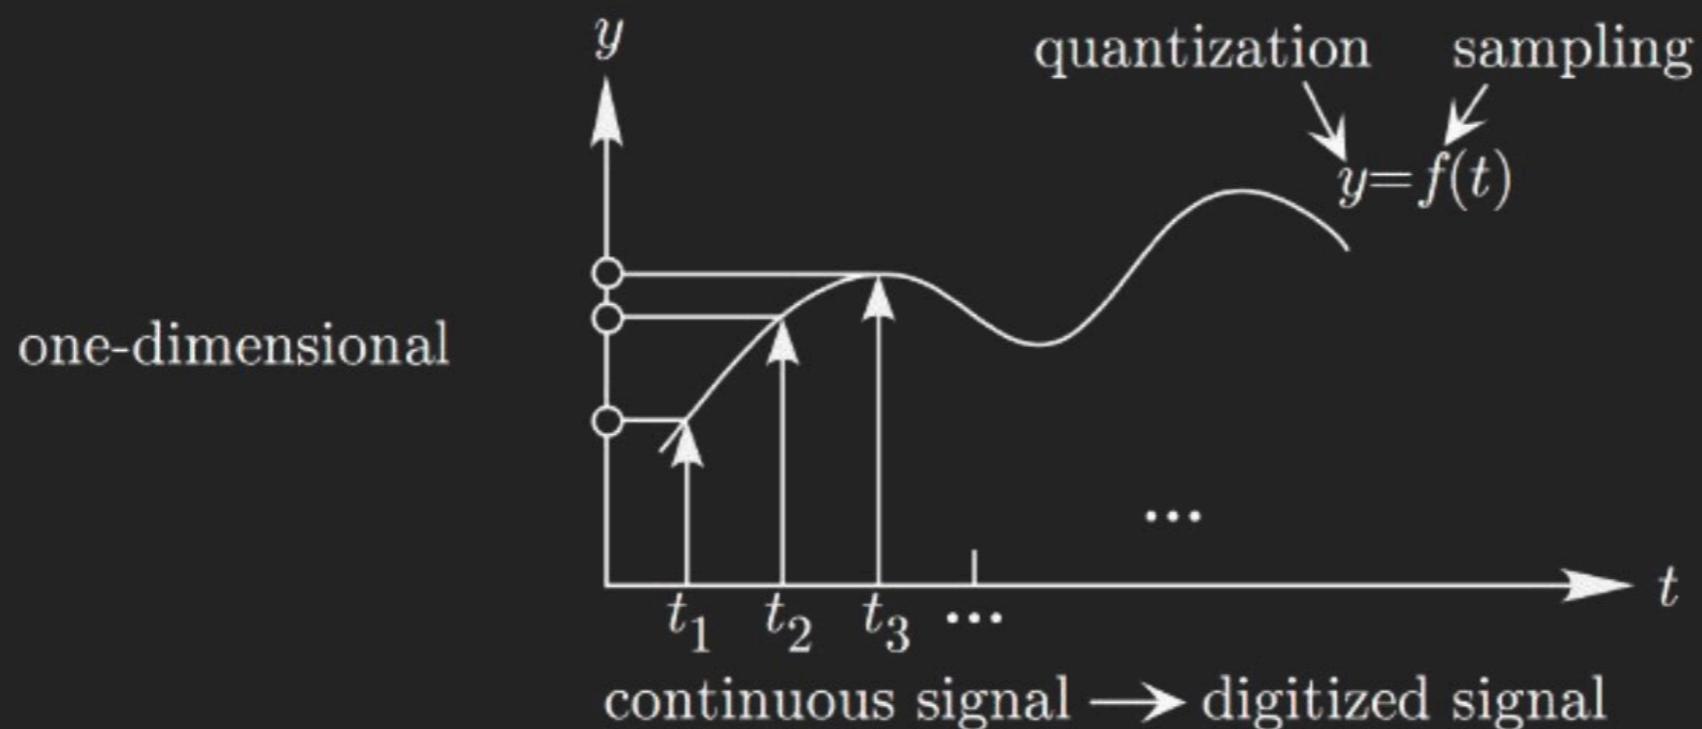
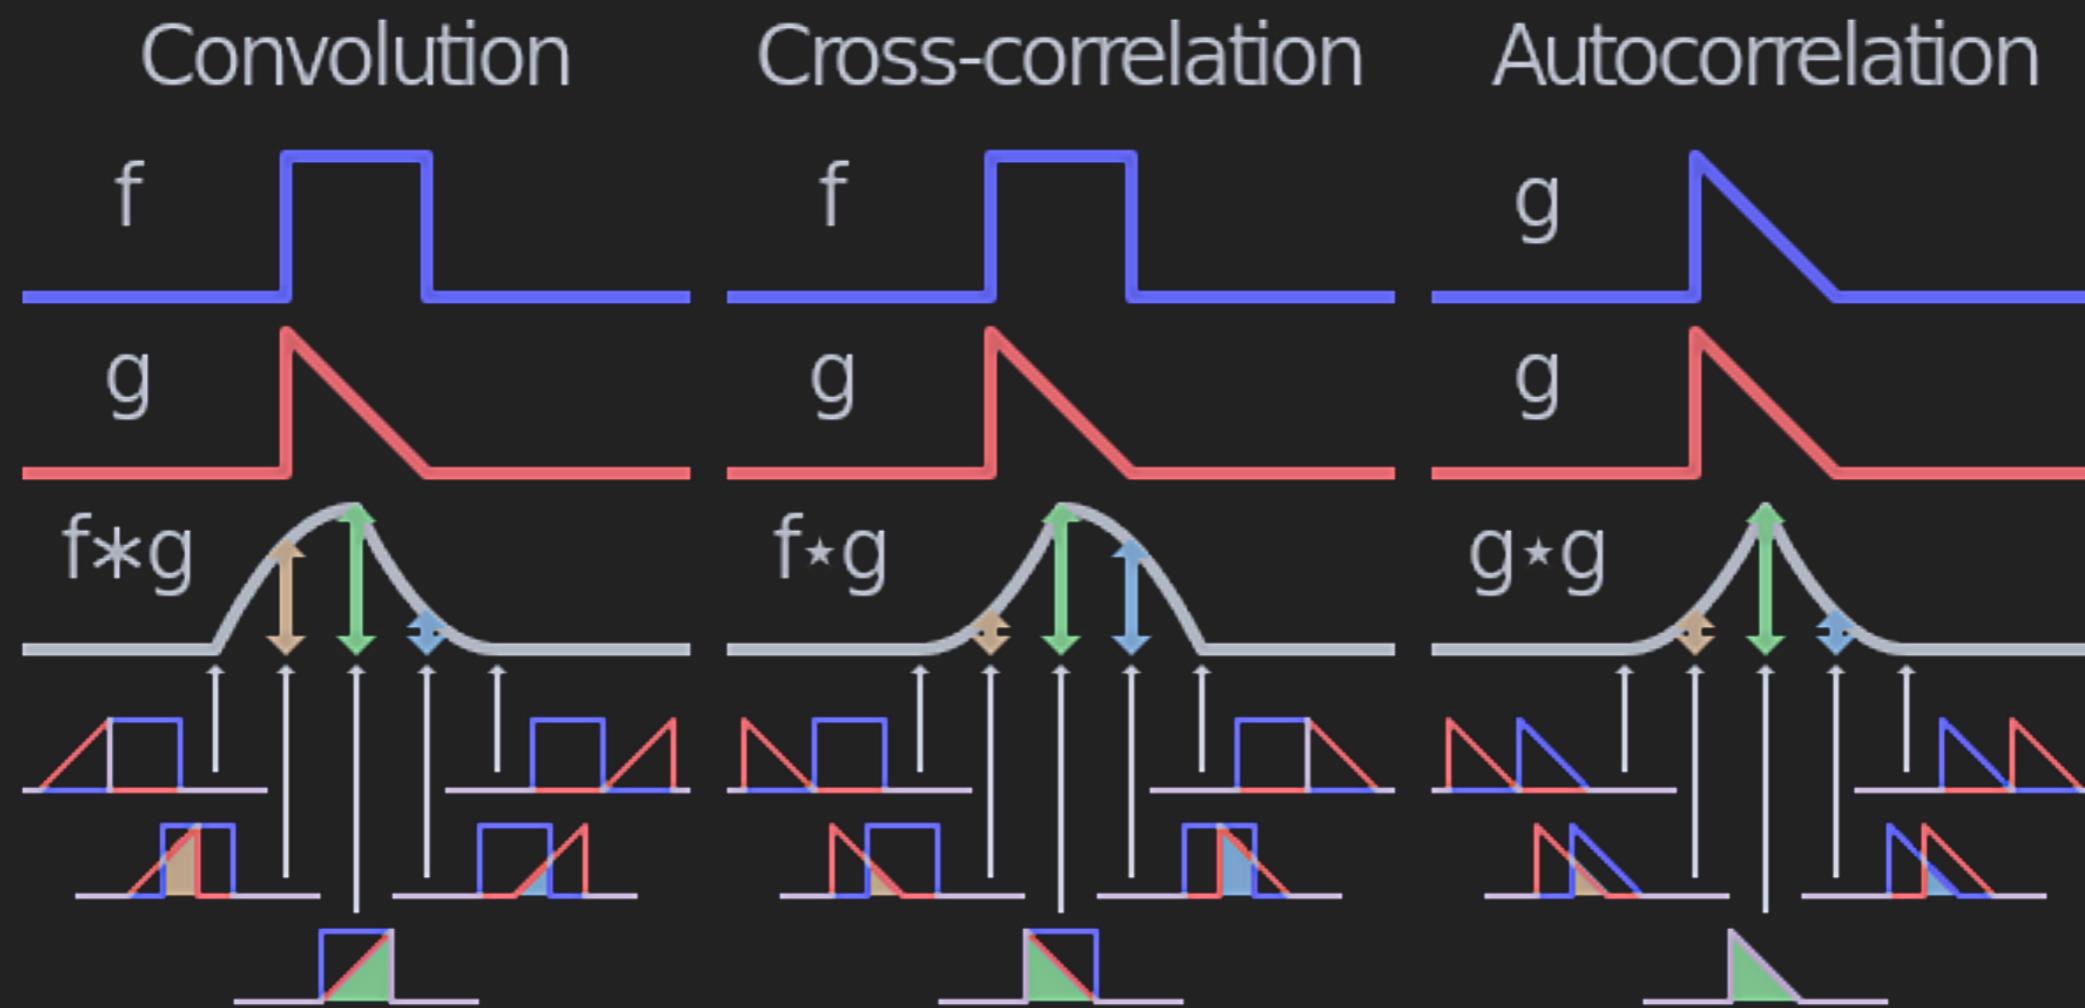


IMAGE FILTERING



CONVOLUTION AND CORRELATION



CORRELATION - LINEAR RELATIONSHIP

$$f \otimes h = \sum_k \sum_l f(k, l)h(l, k)$$

f = Image

h = Kernel

$$\begin{matrix} f \\ \begin{array}{|c|c|c|} \hline f_1 & f_2 & f_3 \\ \hline f_4 & f_5 & f_6 \\ \hline f_7 & f_8 & f_9 \\ \hline \end{array} \end{matrix} \otimes \begin{matrix} h \\ \begin{array}{|c|c|c|} \hline h_1 & h_2 & h_3 \\ \hline h_4 & h_5 & h_6 \\ \hline h_7 & h_8 & h_9 \\ \hline \end{array} \end{matrix} \rightarrow \begin{aligned} f \otimes h = & f_1h_1 + f_2h_2 + f_3h_3 \\ & + f_4h_4 + f_5h_5 + f_6h_6 \\ & + f_7h_7 + f_8h_8 + f_9h_9 \end{aligned}$$

CONVOLUTION

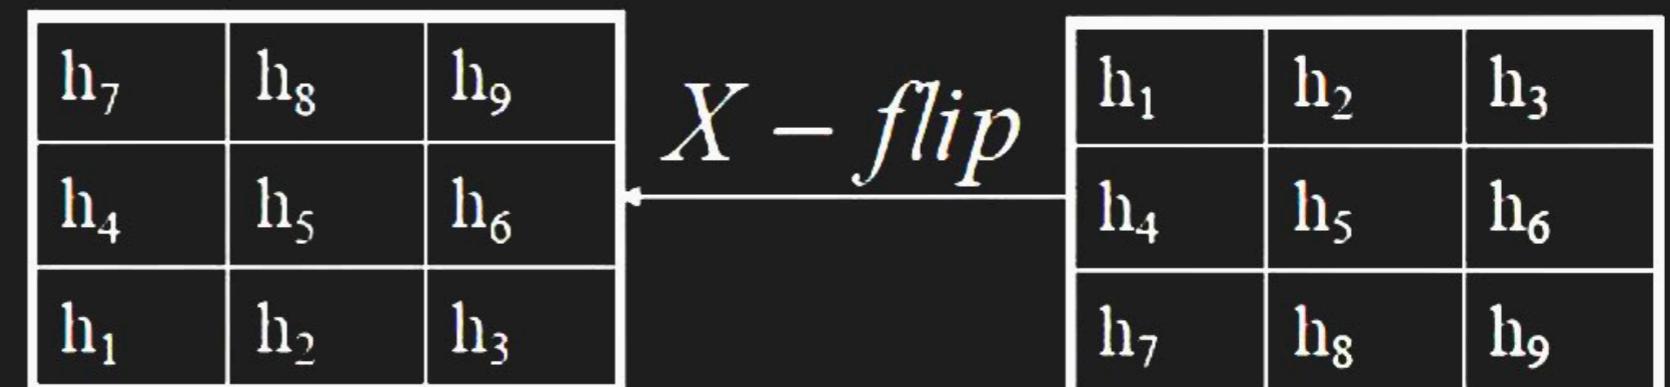
f = Image
 h = Kernel

$$f * h = \sum_k \sum_l f(k, l)h(-k, -l)$$

f_1	f_2	f_3
f_4	f_5	f_6
f_7	f_8	f_9

*

h_9	h_8	h_7
h_6	h_5	h_4
h_3	h_2	h_1



$$\begin{aligned}
 f * h = & f_1 h_9 + f_2 h_8 + f_3 h_7 \\
 & + f_4 h_6 + f_5 h_5 + f_6 h_4 \\
 & + f_7 h_3 + f_8 h_2 + f_9 h_1
 \end{aligned}$$

FILTER = ?

- ▶ Image contains a discrete number of pixels
- ▶ Filtering is forming a new image whose pixel values are a combination of original image's



Not these pixels →

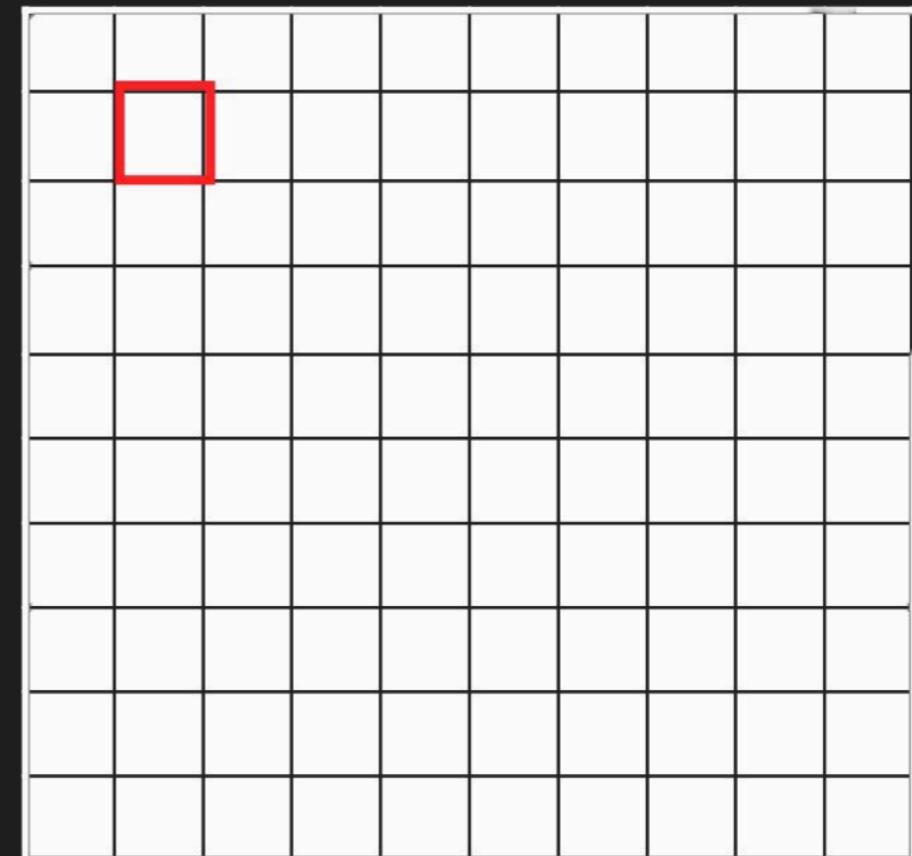
IMAGE FILTERING: MOVING AVERAGE

$$g[\cdot, \cdot] \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$



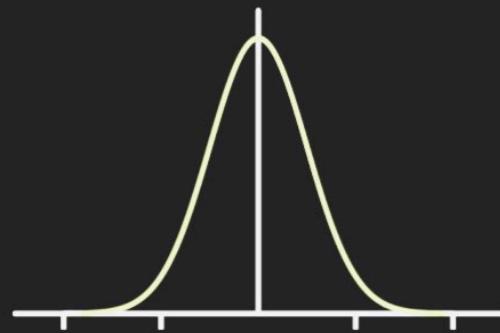
$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

AVERAGE FILTER

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

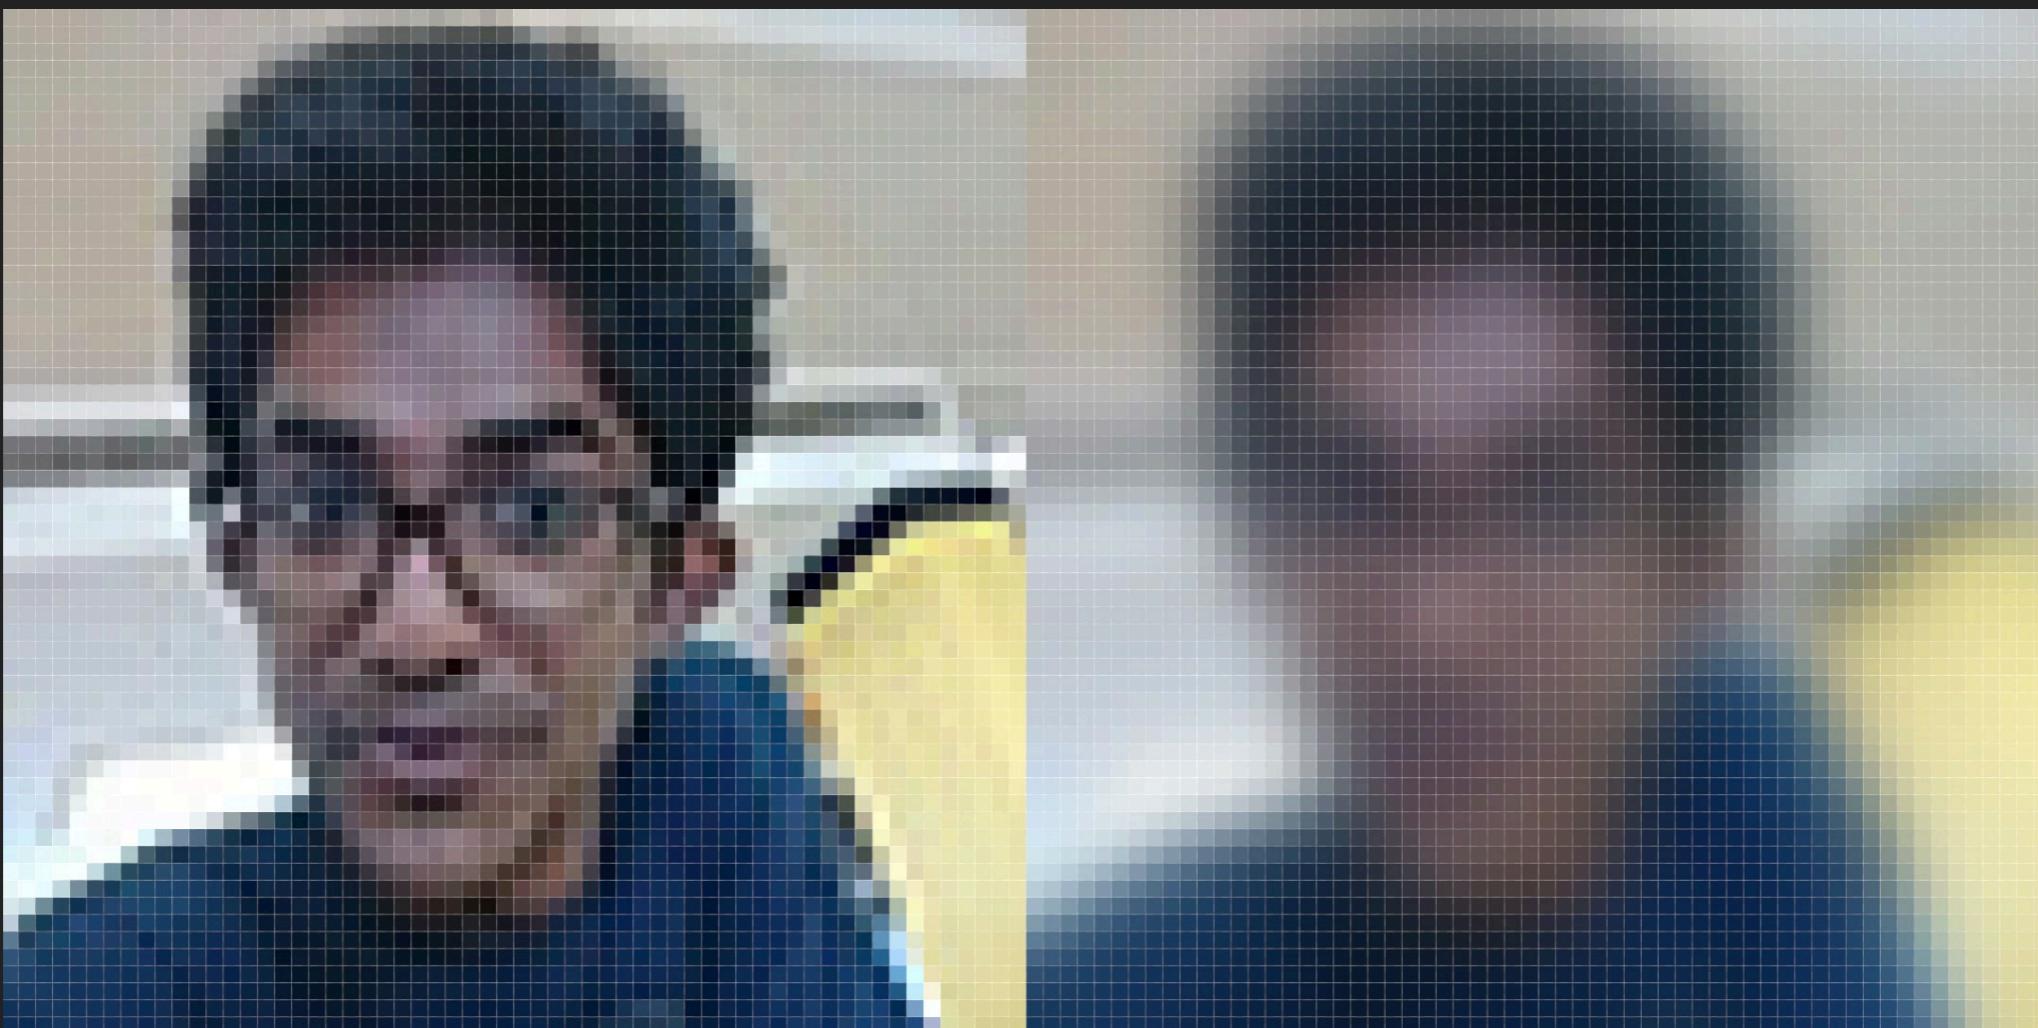


GAUSSIAN FILTER



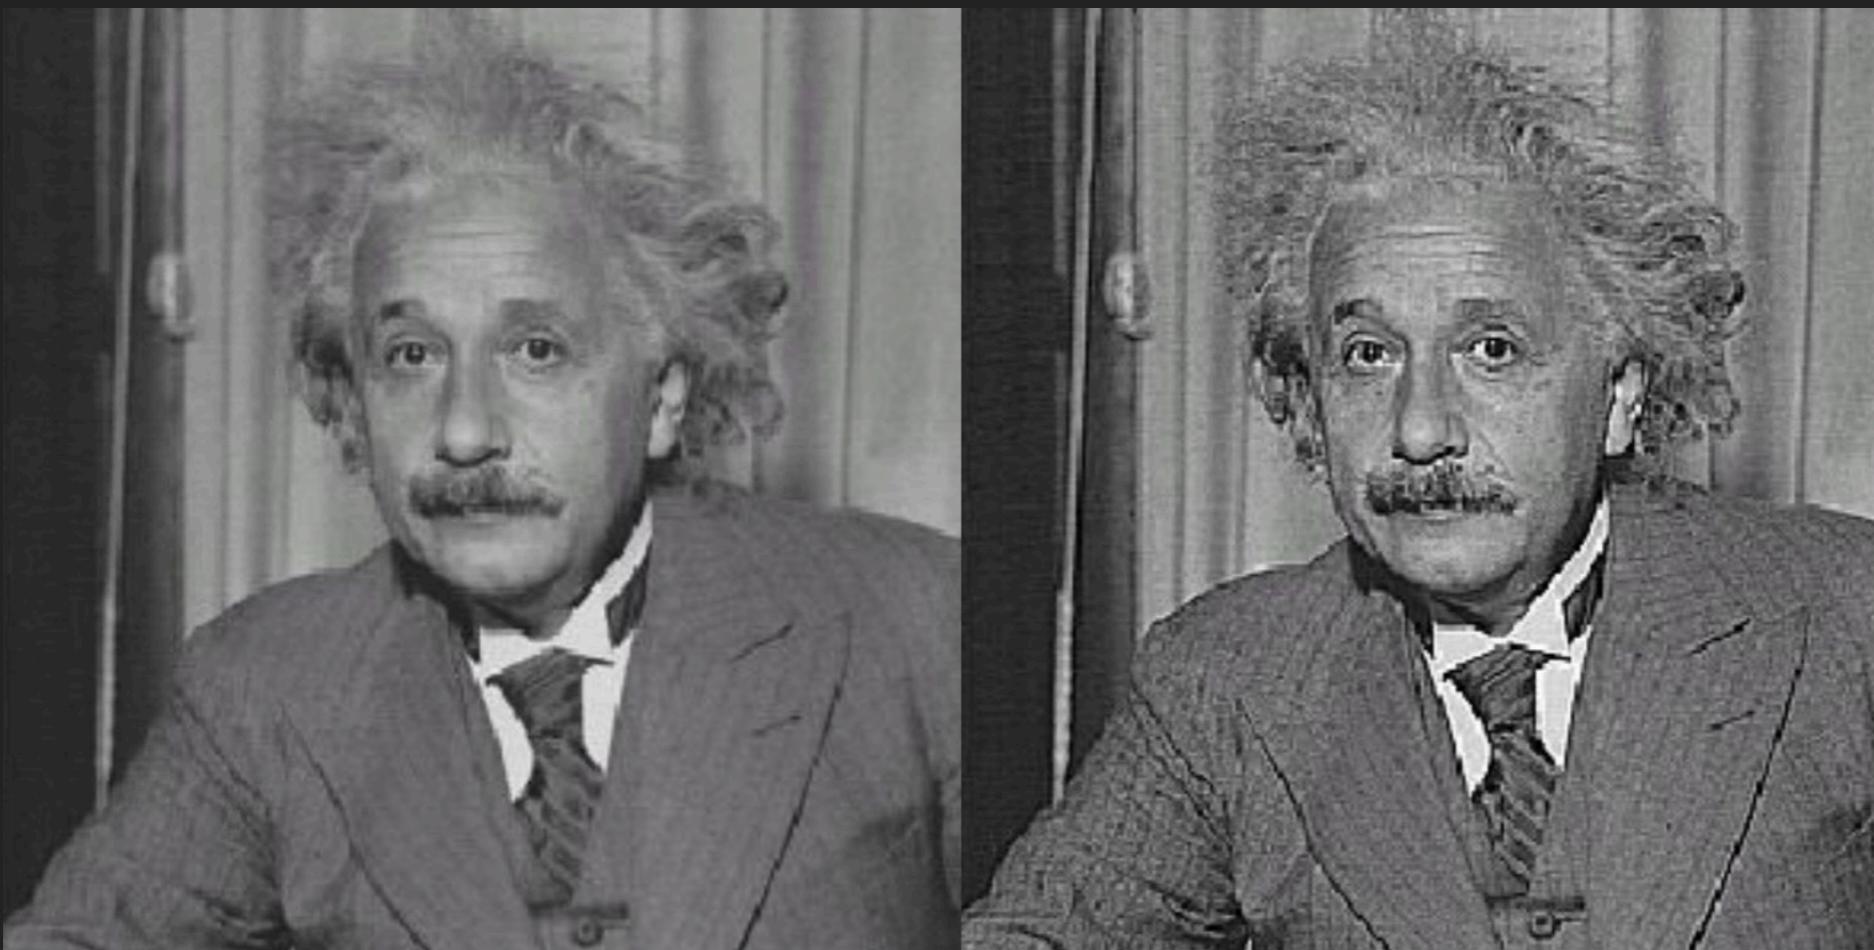
$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

$$\frac{1}{16} \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array}$$



SHARPENING

$$\begin{array}{c|c|c} -1 & -1 & -1 \\ \hline -1 & 9 & -1 \\ \hline -1 & -1 & -1 \end{array}$$



OPERATIONS ON FILTERS



-



=



+



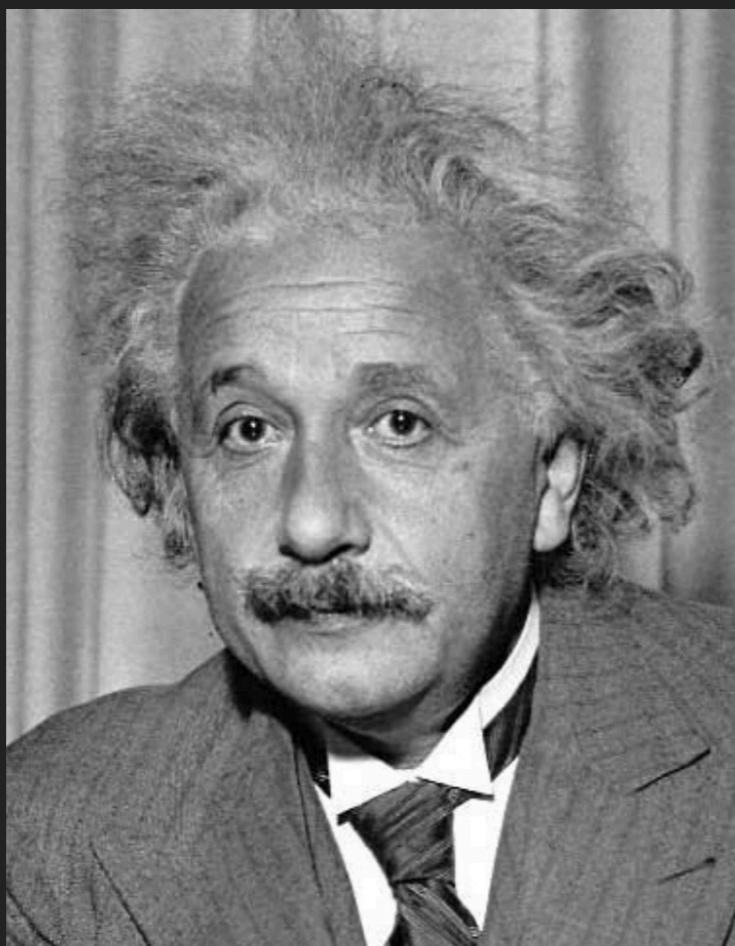
=



SOBEL FILTER

SOBEL X

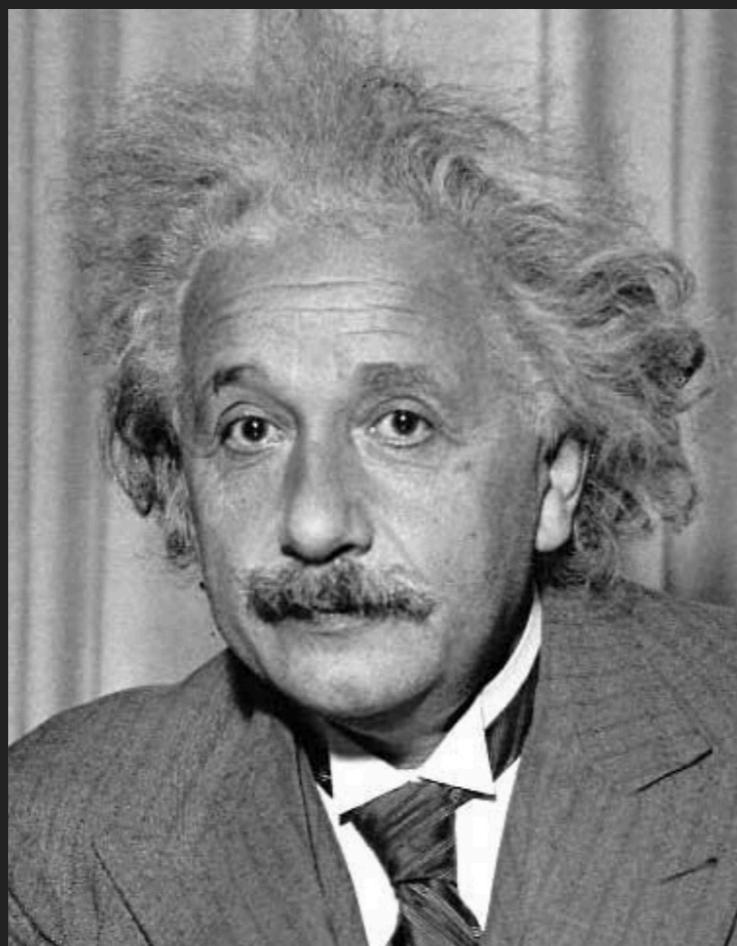
1	2	1
0	0	0
-1	-2	-1



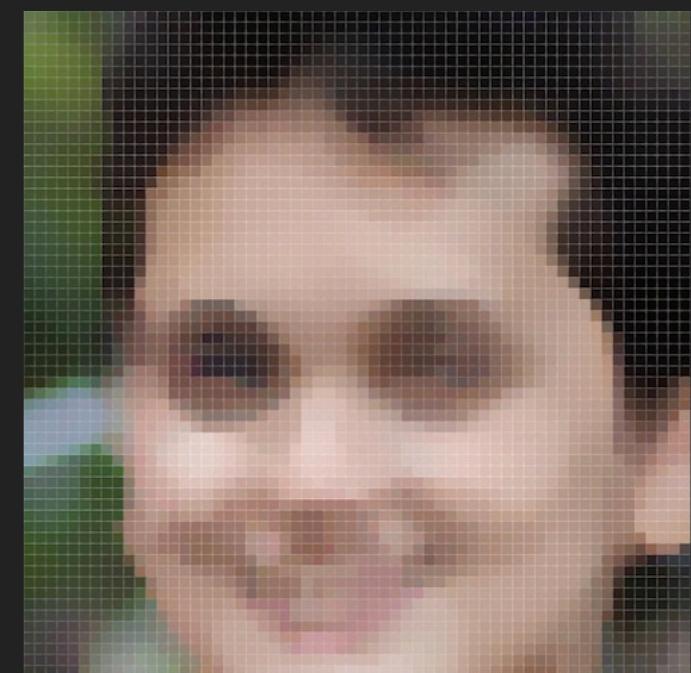
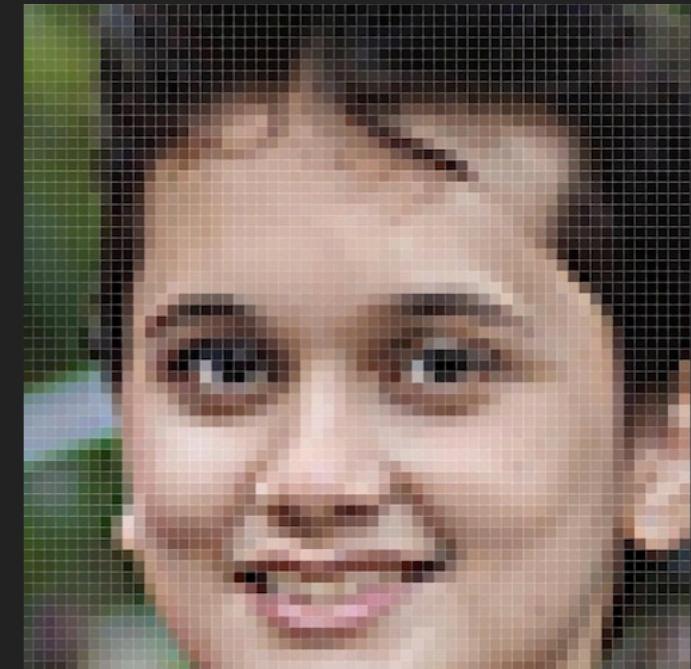
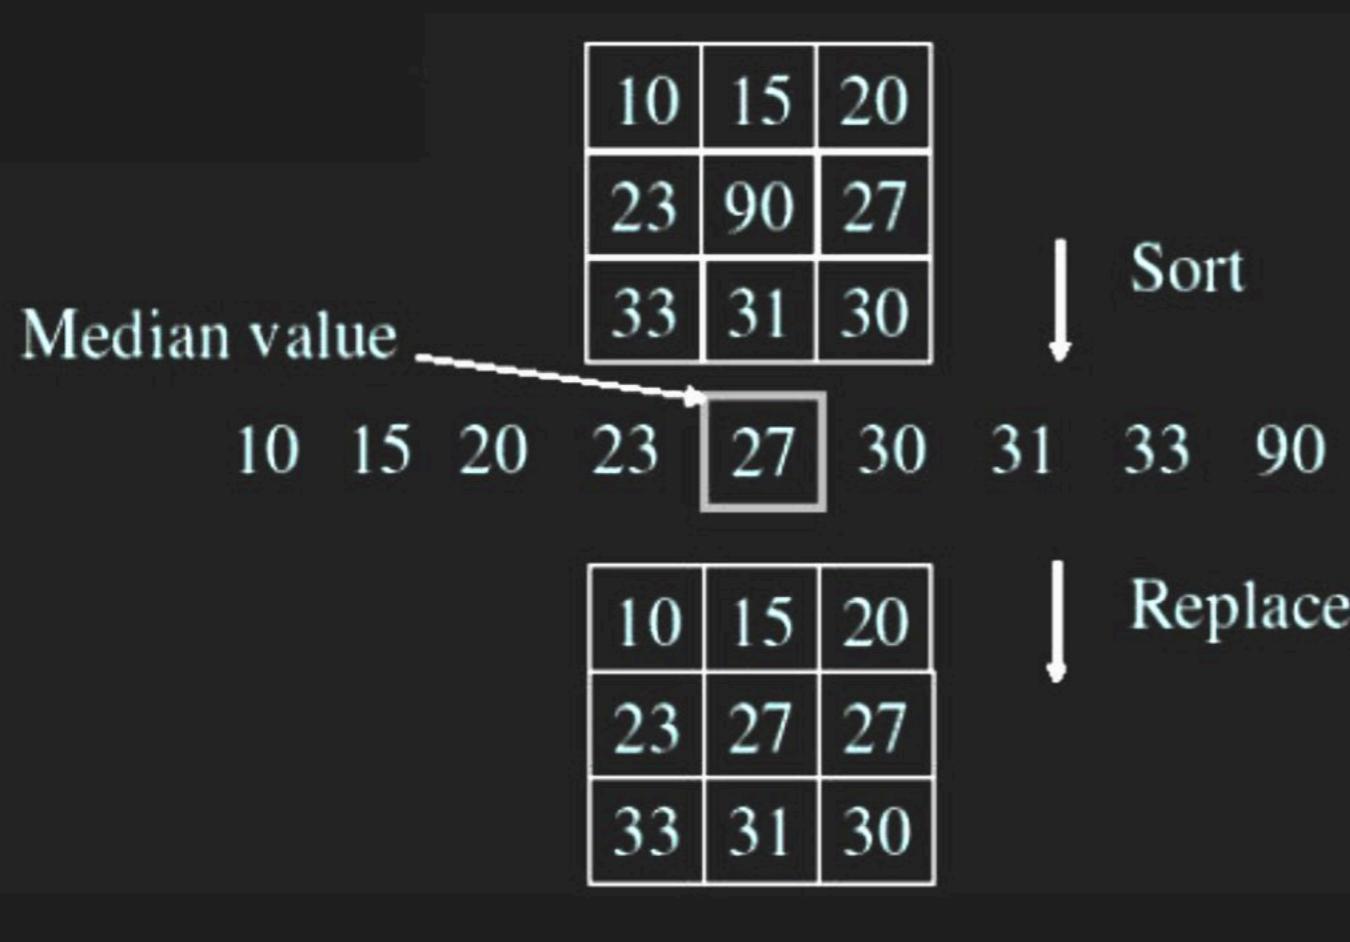
SOBEL FILTER

SOBEL Y

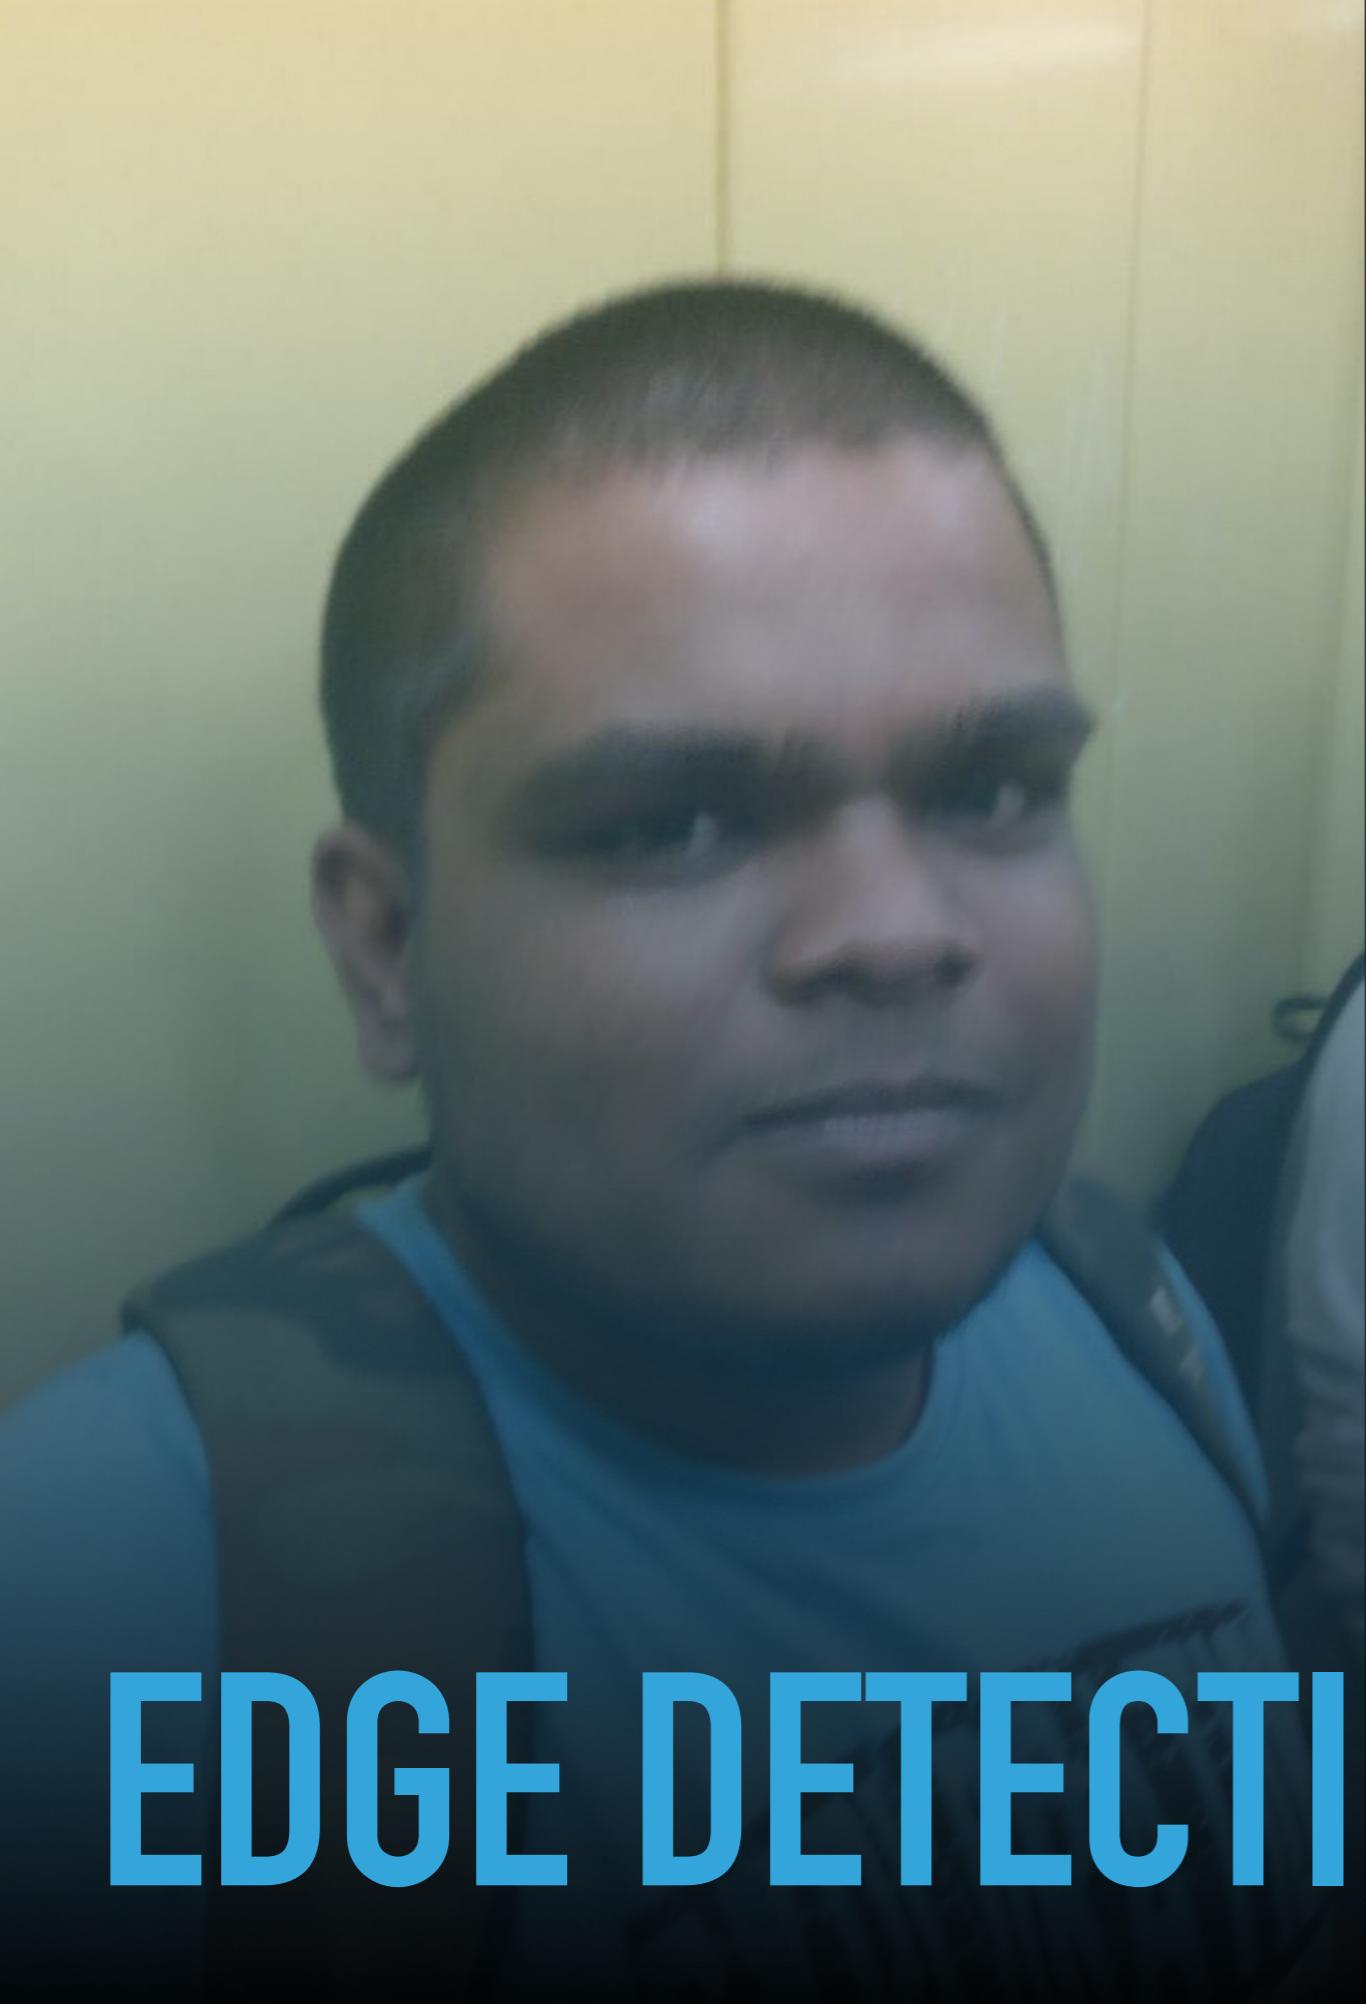
1	0	-1
2	0	-2
1	0	-1



MEDIAN FILTER



EDGE DETECTION



EDGE DETECTION

- ▶ Identify sudden changes (discontinuities) in an image.
- ▶ Intuitively, most semantic and shape information from the image can be encoded in the edge.
- ▶ Edges extract information and recognize edges
- ▶ Recover geometry and viewpoint

WHAT EXACTLY IS AN EDGE

An edge is a place of rapid change in the image intensity function

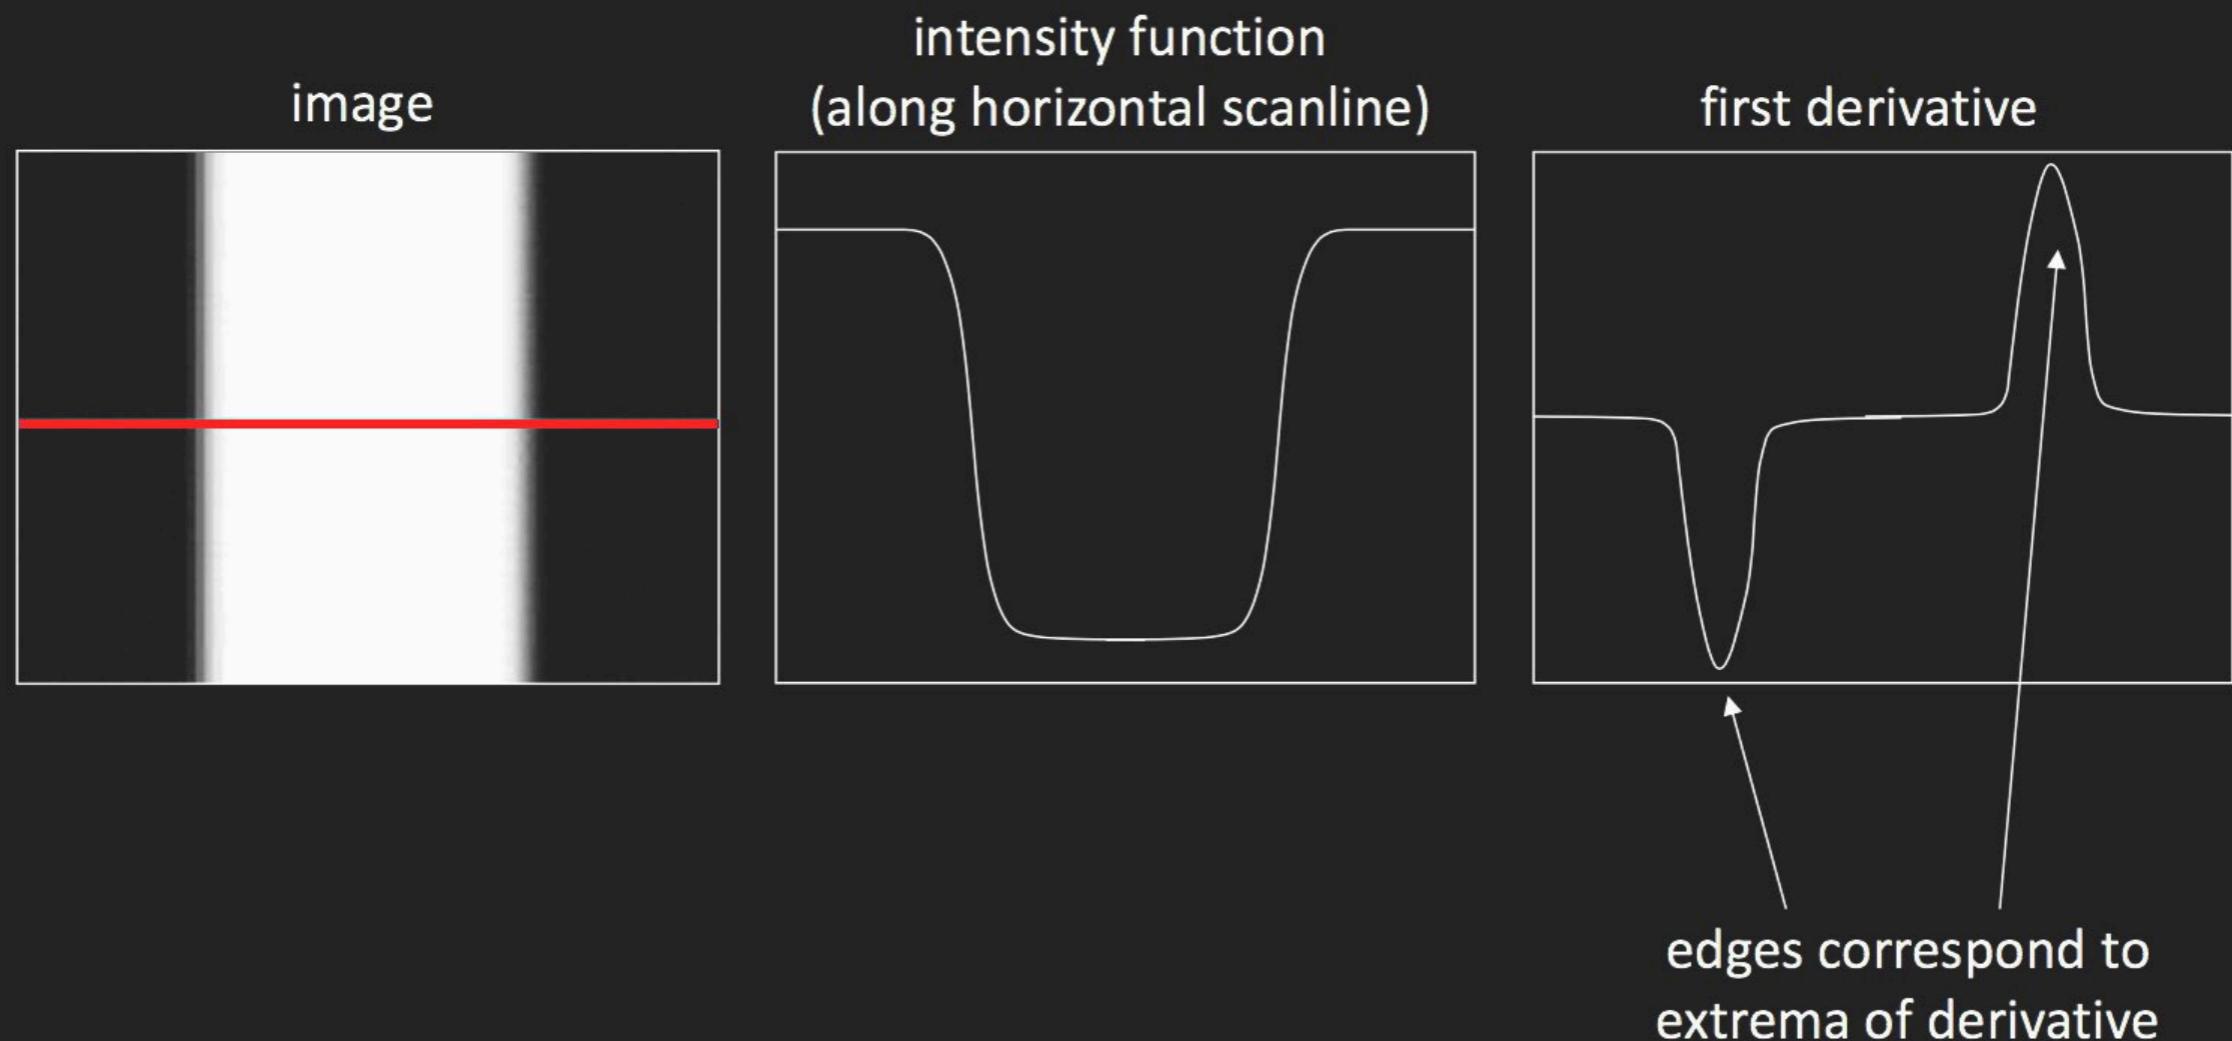
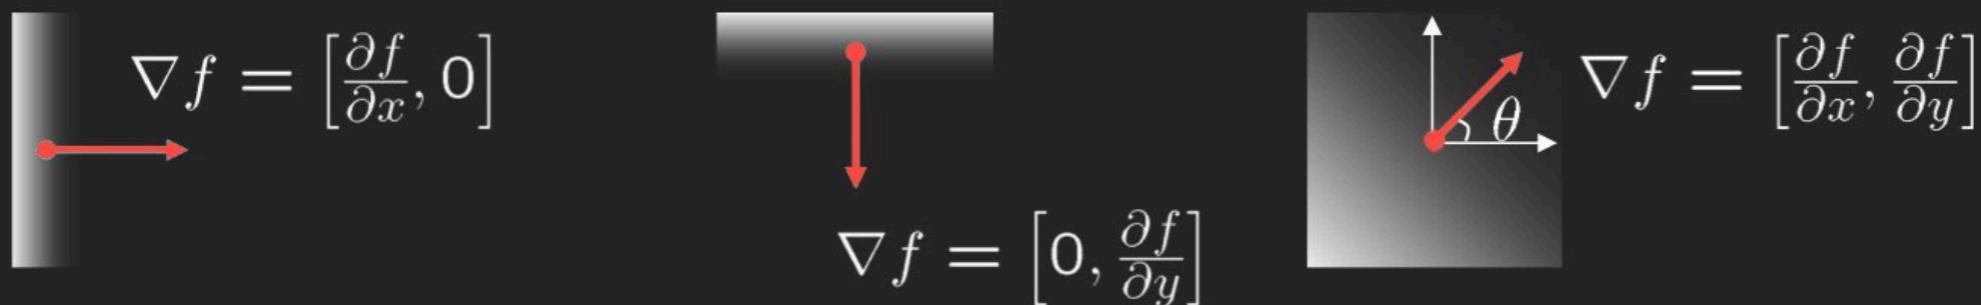


IMAGE GRADIENT

- The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$



The gradient points in the direction of most rapid increase in intensity

- How does this direction relate to the direction of the edge?

The gradient direction is given by

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$

Source: Steve Seitz

CANNY EDGE DETECTOR

- ▶ Smooth image with a *Gaussian*
 - ▶ optimizes the trade-off between noise filtering and edge localization
- ▶ Compute the *Gradient* magnitude using approximations of partial derivatives. (2x2 filters)
- ▶ *Thin* edges by applying non-maxima suppression to the gradient magnitude.
- ▶ Detect edges by *double thresholding*

GRADIENT

- ▶ Convolve with G_x and G_y at each point.
- ▶ Magnitude and gradients are computed

$$M[i, j] = \sqrt{P[i, j]^2 + Q[i, j]^2}$$

$$\theta[i, j] = \tan^{-1}(Q[i, j], P[i, j])$$

$$G_x = \begin{bmatrix} -1 & 1 \\ -1 & 1 \end{bmatrix}$$

$$G_y = \begin{bmatrix} 1 & 1 \\ -1 & -1 \end{bmatrix}$$

- ▶ Floats are dropped for faster computation

NON MAXIMA SUPPRESSION

- ▶ Thin edges by keeping values of gradient
 - ▶ Not always at the location of the edge
 - ▶ There are many thick edges

GRADIENT ORIENTATION

- ▶ Get the four basic directions of the edges orientation: horizontal, vertical, 45° and -45° .
- ▶ Find the direction closest to $a(x, y)$.
- ▶ If the value of $M(x, y)$ is less than at least one of its two neighbors along d_k , let $g_N(x, y) = 0$ (suppression); otherwise, let $g_N(x, y) = M(x, y)$

THRESHOLDING

- ▶ Final operation is to threshold $g_N(x, y)$ to reduce false edge points.

Hysteresis thresholding:

$$g_{NH}(x, y) = g_N(x, y) \geq T_H$$

$$g_{NL}(x, y) = g_N(x, y) \geq T_L$$

and

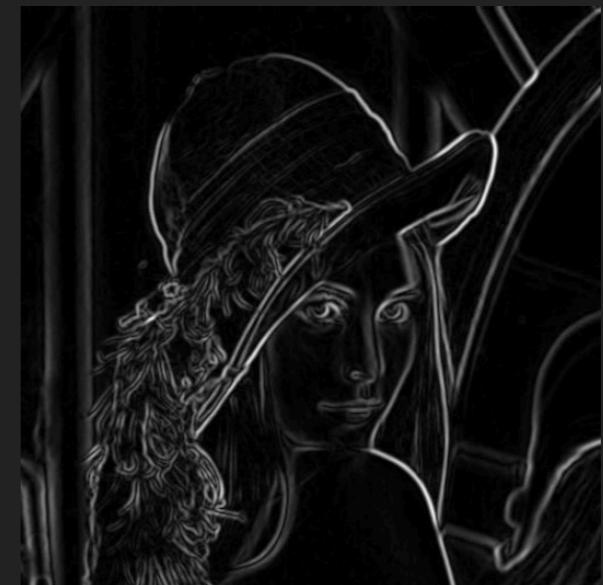
$$g_{NL}(x, y) = g_{NL}(x, y) - g_{NH}(x, y)$$



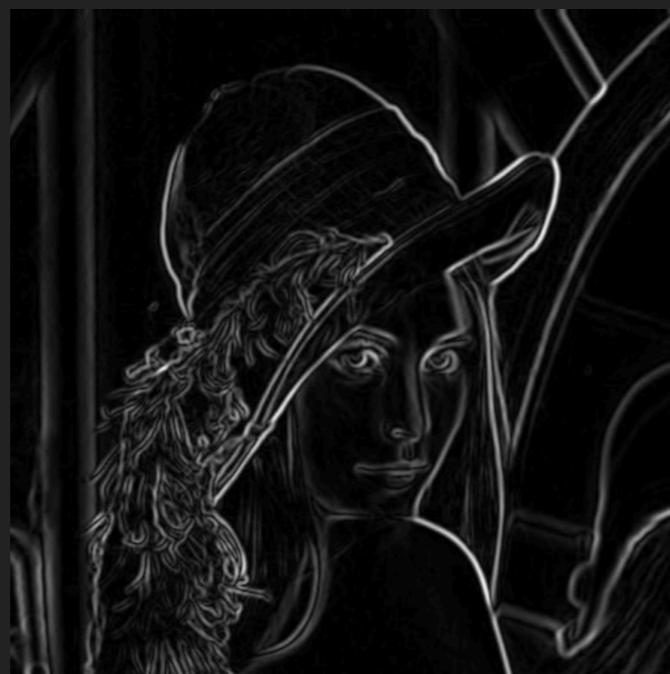
gradient X



gradient Y



gradient magnitude



Before non maxima
suppression



After non maxima
suppression

CANNY EDGE DETECTOR





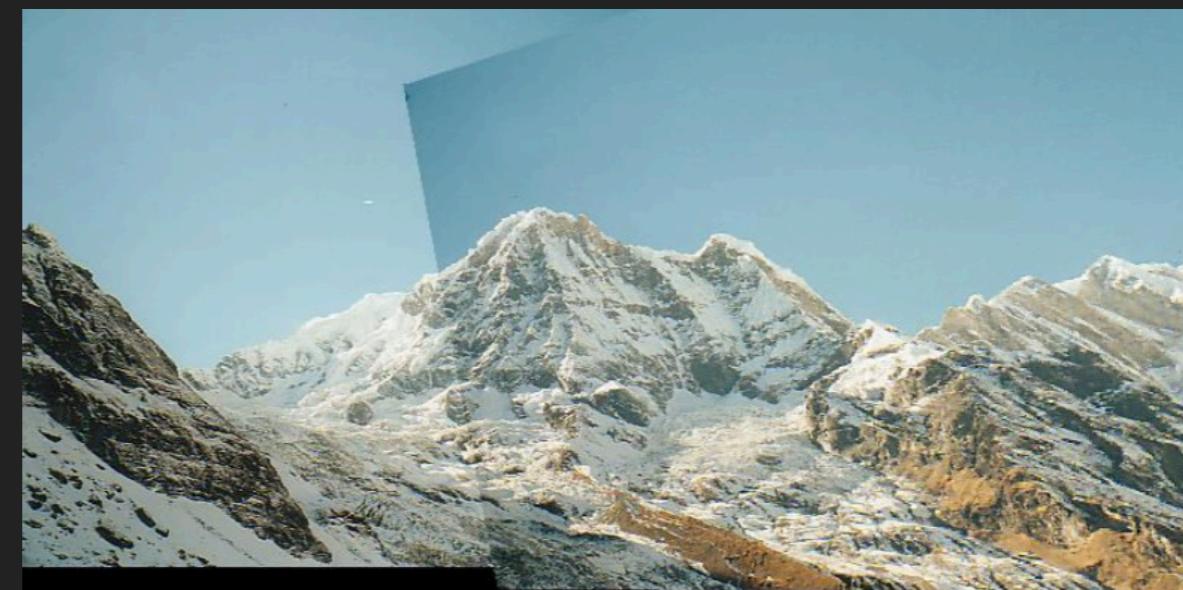
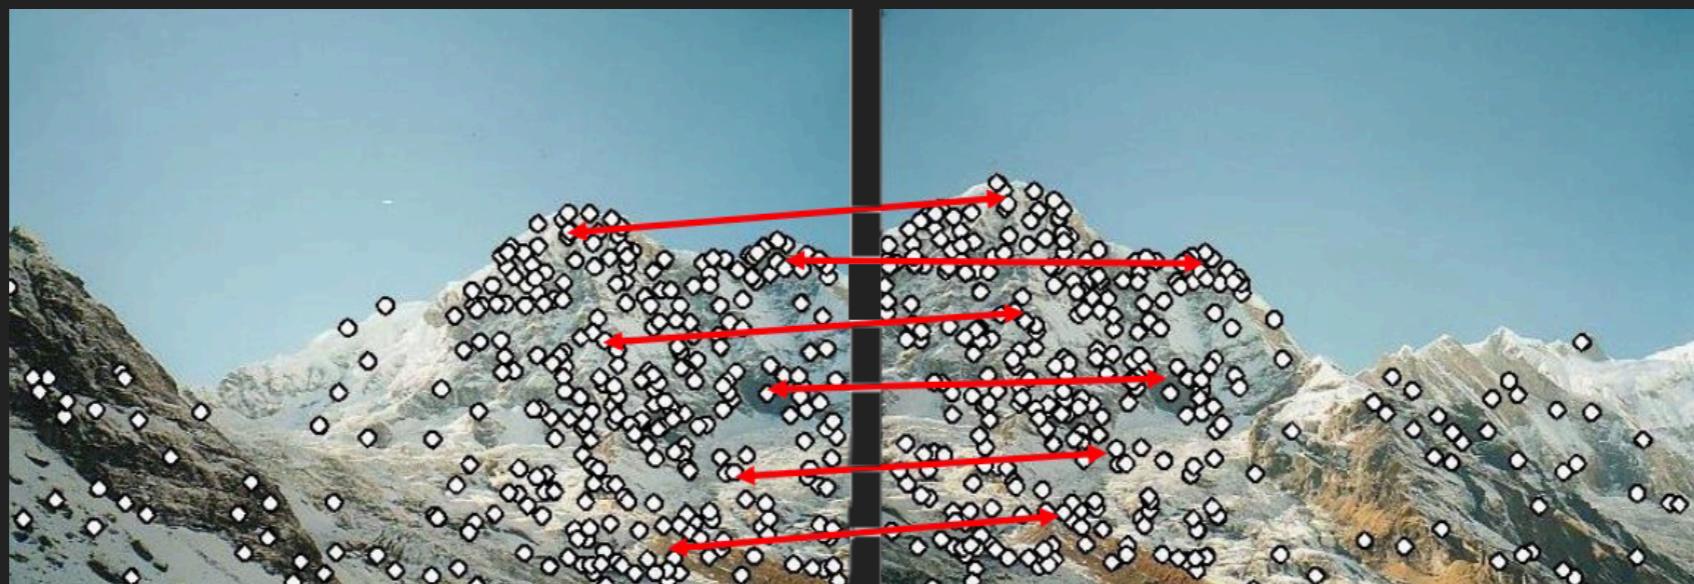
FEATURE DETECTION

GENERAL STEPS

- ▶ Find a set of distinctive key points
- ▶ Define a region around each key point
- ▶ Extract and normalize the region content
- ▶ Compute a local descriptor from a normalized region
- ▶ Match local descriptors

IMAGE STITCHING

- ▶ Extract features, combine feature, and align image



HARRIS DETECTOR FORMULATION

Change of intensity for the shift $[u,v]$:

$$E(u,v) = \sum w(x,y) [I(x+u, y+v) - I(x, y)]^2$$

The diagram illustrates the Harris detector formulation. It shows the formula $E(u,v) = \sum w(x,y) [I(x+u, y+v) - I(x, y)]^2$. Three components are highlighted with arrows pointing to them: 'Window function' points to the term $w(x,y)$; 'Shifted intensity' points to the term $I(x+u, y+v)$; and 'Intensity' points to the term $I(x, y)$.

- ▶ Window function $w(x, y)$ can be linear or gaussian

This measure of change can be approximated by:

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a 2×2 matrix computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

↑
Sum over image region – the area we are
checking for corner

**Gradient with
respect to x ,
times gradient
with respect to y**

$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y]$$

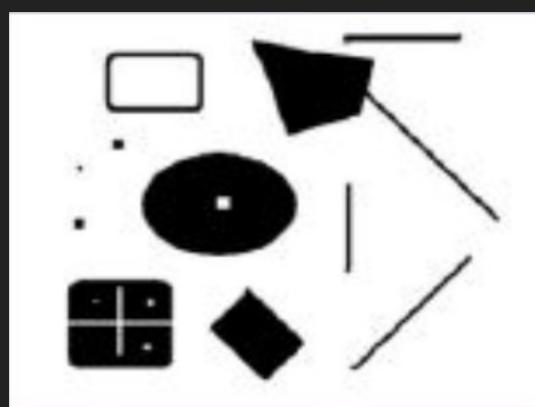
where M is a 2×2 matrix computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

↑
Sum over image region – the area we are
checking for corner

**Gradient with
respect to x ,
times gradient
with respect to y**

$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y]$$



I



Ix

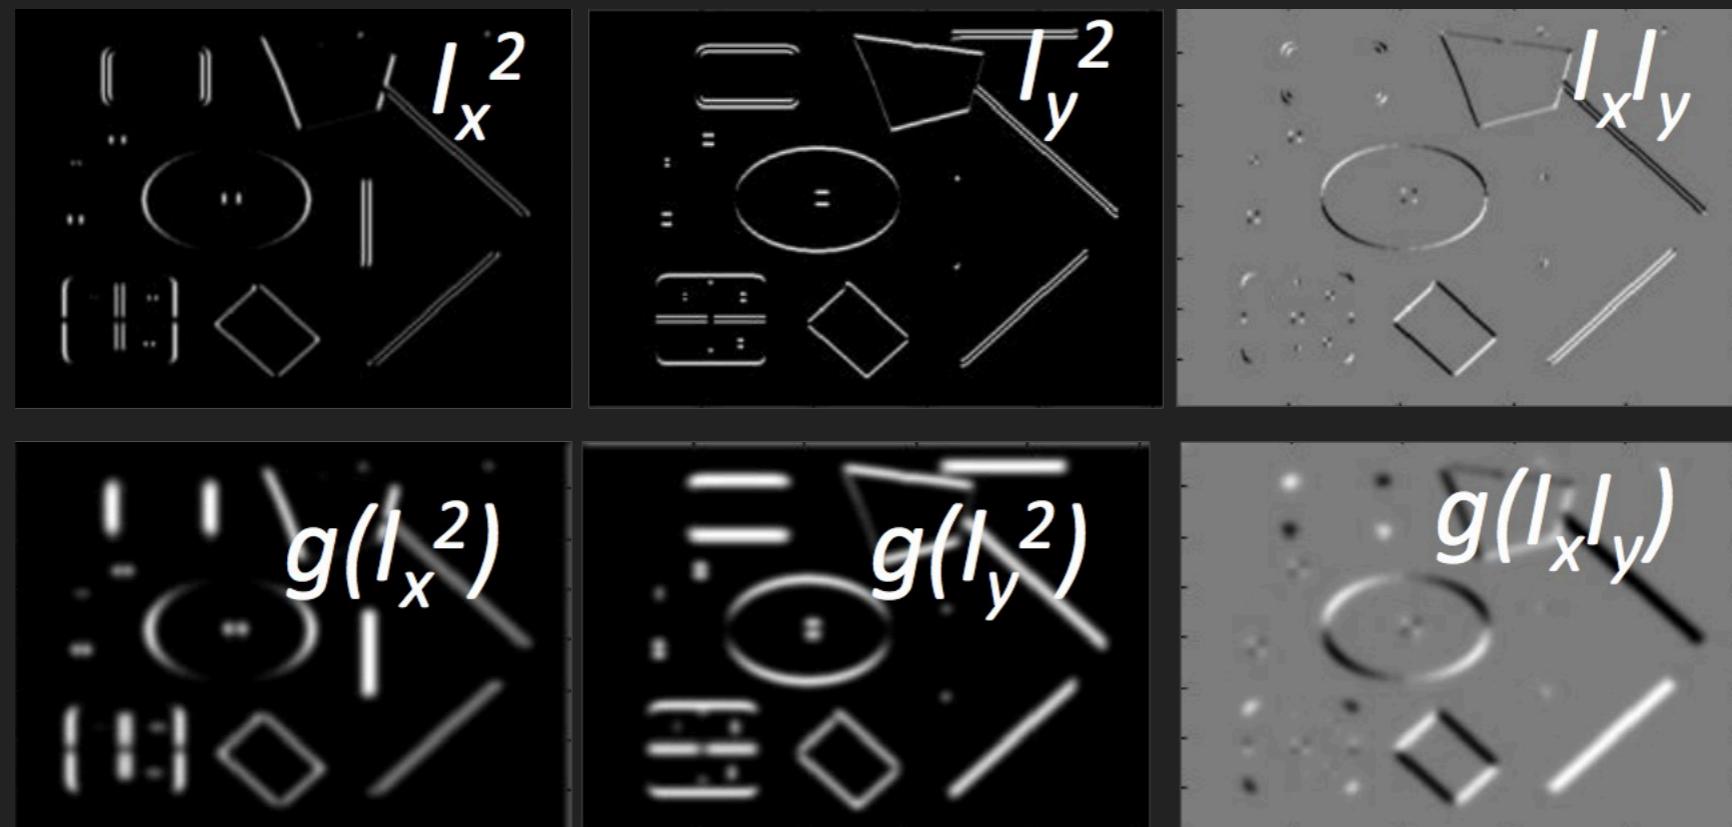


Iy



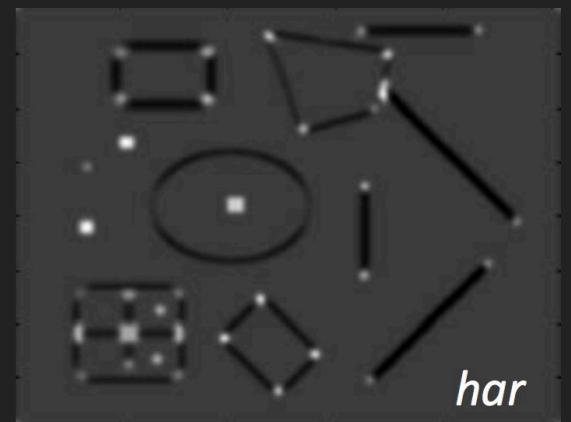
Ix Iy

- ▶ Gaussian filter $g(\sigma)$ on I_x^2 , I_y^2 , and $I_x I_y$.



- ▶ Cornerness function – both eigenvalues are strong

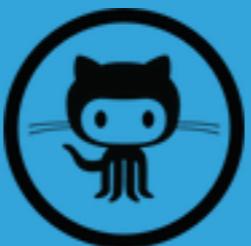
$$\begin{aligned}
 har &= \det[M(\sigma_I, \sigma_D)] - \alpha[\text{trace}(M(\sigma_I, \sigma_D))^2] \\
 &= g(I_x^2)g(I_y^2) - [g(I_x I_y)]^2 - \alpha[g(I_x^2) + g(I_y^2)]^2
 \end{aligned}$$







THANK YOU!



@sidkothiyal
@avikantz